

INTRODUZIONE AI BIG DATA

Prof. Flavio Venturini
Lez.03 - Introduzione ad Hadoop

INTRODUZIONE AD HADOOP

Topics

- Intro to MapReduce
- HDFS architecture
- Accessing HDFS

INTRO TO MAPREDUCE

The problem

It's all about storage and transfer rates

1990: 1.370 MB – 4.4MB/s
5 mins to read it

2015: 1 TB – 100MB/s
2.5 hrs to read it

Possible solution

Obvious: read multiple disks in parallel

Issues:

- Share of resources
- Hardware failures
- Combination of results

Answer to issues



a **reliable**, **scalable** platform for shared storage and analysis, which is open source and runs on commodity hardware, i.e. **affordable**




UNIVERSITÀ TELEMATICA
INTERNAZIONALE UNINETTUNO

Copyright © Università Telematica Internazionale UNINETTUNO

INTRODUZIONE AI BIG DATA

Prof. Flavio Venturini
Lez.03 - Introduzione ad Hadoop

Basic Hadoop Architecture



Example: website logs

We want to identify the behaviour of the users of our website:

time page impression
number page impression
etc...

Input data format

\$6 \$11 \$15

date time s-ip cs-method
cs-uri-stem cs-username cs-port
cs-uri-info c-ip cs(User-Agent)
cs(Page) sc-status sc-substatus
sc-win32-status time-taken

Input data example

2016-01-27 09:11:00 10.130.232.47 POST
/_vti_bin/sites.asmx filplocco 443 cpi.bisceglie
93.63.43.25 Mozilla/5.0+(compatible;+MSIE+9.0;
+Windows+NT+6.0;+Trident/5.0)
[https://www.mysite.it/AreaRiservata/Operatori/
Garanzie/Pagine/Profiling.aspx](https://www.mysite.it/AreaRiservata/Operatori/Garanzie/Pagine/Profiling.aspx) 200 0 0 281

Analysing with Unix tools

```
awk '{print $6," ", $15}' file.log |  
awk '{user[$1]++} END {for (var  
in user) print var," ",user[var]}
```

Shows the number of pages
visited by each named user:

filplocco 33

Exercise

Using awk, write a script to
sum the total time spent by
each user on the website

Hint: See previous example.
You have already extracted the
two variables you need



INTRODUZIONE AI BIG DATA

Prof. Flavio Venturini
Lez.03 - Introduzione ad Hadoop

Analyzing data with Hadoop

MapReduce job:


Input data, MapReduce program,
Configuration info

Input split:


It is the chunk of input data given
to a map task for processing

Analyzing data with Hadoop

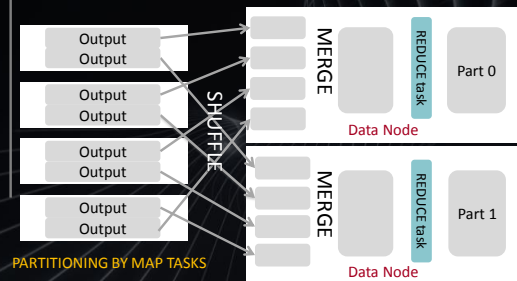
1: Map tasks



2: Reduce (SINGLE)



2: Reduce (MULTIPLE)



Reduce details

The number of Reduce tasks is
not governed by the size
of the input, but is specified
independently

Example Map Reduce

Problem:

I want to find the maximum
time spent on each visited
page of my website



INTRODUZIONE AI BIG DATA

Prof. Flavio Venturini
Lez.03 - Introduzione ad Hadoop

Example Python Map task

```
import sys
for line in sys.stdin:
    res = line.split()
    print (res[10],res[14])
```

This creates a data file with -on each line- the visited page and the time spent on it

Ex. Python Reduce task

```
import sys
(last_key, max_val)=(None, -sys.maxint)
for line in sys.stdin:
    (key, val) = line.split()
    if last_key and last_key !=key:
        print (last_key, max_val)
        (last_key, max_val)=(key, int(val))
    else:
        (last_key,max_val)=(key,max(max_val,int(val)))
if last_key:
    print (last_key, max_val)
```


Review Questions

- What are the limits of working on very large files using the awk-like paradigm?
- Describe the MapReduce paradigm

Review Questions

- When is there a transfer of data between nodes in a MapReduce operation?

HDFS ARCHITECTURE



HDFS Design

It is a file system designed for storing **very large files** with **streaming data access** patterns, running on **clusters of commodity hardware**



INTRODUZIONE AI BIG DATA

Prof. Flavio Venturini
Lez.03 - Introduzione ad Hadoop

The block concept

A disk has a block size which is the minimum amount of data that it can read or write

Normally
512 bytes

Filesystems work on multiple disk blocks

Few KB

The HDFS Block

HDFS block is 128MB by default

Benefits:

- a file can be larger than any single disk on the network
- Simplification of the storage management

Name and Data nodes

An HDFS cluster is made of: one or more **name nodes** (masters) and a set of **data nodes** (workers)

Name nodes maintain the filesystem namespace

Datanodes store and retrieve blocks

Block caching

Datanode's memory can be used to cache block

A given block is cached only in one datanode

Example: look-up table used in a join

HDFS Federation

Namenode memory is limited, so it limits scalability

HDFS federation allows each namenode to manage a portion of the filesystem namespace (**namespace volume** and **block pool**)

HDFS and HA

Replication of namenode metadata + secondary namenodes = data loss prevention

but
it is not High Availability

NAMENODE IS A SINGLE POINT OF FAILURE

Hadoop 2 introduced NN active-standby



INTRODUZIONE AI BIG DATA

Prof. Flavio Venturini
Lez.03 - Introduzione ad Hadoop

Review Questions

- Describe the difference between a Namenode and a Datanode
- Why the loss of a DataNode is not an issue?
- Why is it an issue the loss of a namenode if not properly managed?

ACCESSING HDFS

Using Filesystem APIs

FileSystem is an abstract class that represents a generic file system. To create an instance of the HDFS, you call the method **FileSystem.get()**. Then **HdfsWriter** class calls the **create()** method to create a file in HDFS

Using Filesystem APIs

HdfsReader calls the method **open()** to open a file in HDFS, which returns an **InputStream** object that can be used to read the contents of the file

Example

Read file from the local file system and write it to HDFS

```
Configuration conf = getConf();  
OutputStream os = fs.create(outputPath);  
InputStream is = new BufferedInputStream(new  
FileInputStream(localInputPath));  
IOUtils.copyBytes(is, os, conf);
```

Directories

FileSystem provides method **mkdirs(Path f)** to create them. It creates all the path if it does not exist

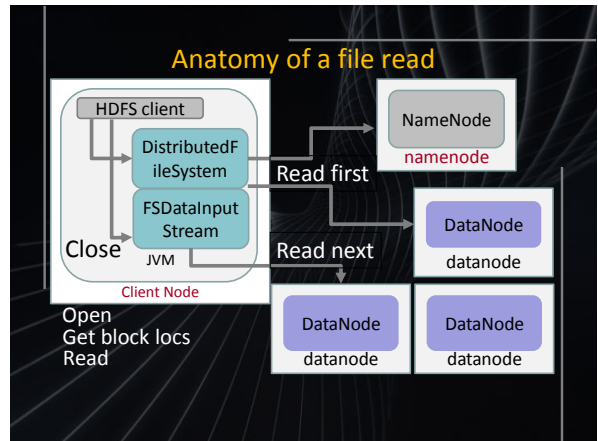


INTRODUZIONE AI BIG DATA

Prof. Flavio Venturini
Lez.03 - Introduzione ad Hadoop

Filesystem operations

Listing files: `listStatus`
File info: `getFileStatus`
File patterns: `globStatus`
Deletion: `delete`



Review Questions

- What are the basic classes to be used for accessing HDFS with FileSystem?
- Explain how Data Read process works.

SUMMARY QUESTIONS

- What are the limits of working on very large files using the awk-like paradigm?
- Describe the MapReduce paradigm
- When is there a transfer of data between nodes in a MapReduce operation?

- Describe the difference between a Namenode and a Datanode
- Why the loss of a DataNode is not an issue?
- Why is it an issue the loss of a namenode if not properly managed?



INTRODUZIONE AI BIG DATA

Prof. Flavio Venturini
Lez.03 - Introduzione ad Hadoop

- What are the basic classes to be used for accessing HDFS with FileSystem?
- Explain how Data Read process works.

